

COMP3620 / 6320 – S1 2023

程序代写代做 CS 编程辅导



This lecture will be recorded

Turn off your camera if you do not want to be in the recording
WeChat: cstutorcs

If you cannot see my whole slide, go to View Options -> Zoom Ratio and select
~~Assignment Project Exam Help~~ Fit to Window

If you follow the class online and want to ask a question, either:

- Raise your hand and I will open ~~QQ 749389476~~, or
- Type the question in the chat box <https://tutorcs.com>
- Since the lecture is live in person, I cannot constantly monitor the online part, so please be patient

Reminders

程序代写代做 CS 编程辅导



- **Class representatives**

- Please nominate yourself ~~by sending a private message on piazza by today.~~
- You are free to nominate yourself whether you are currently on-campus or studying remotely. WeChat: cstutorcs

- **Labs and Tutorials**

Assignment Project Exam Help

- Sign-up available on wattle. Please use ~~the same~~ ~~E-mail tutorcs@anu.edu.au~~ on Mytimetable, wattle selection has preference. QQ: 749389476
- No need to sign up for labs
- First tutorial and lab: ~~14-20 March 2023 (Tuesday to Monday)~~ <https://tutpros.com>
- **First tutorial question sheet is out now on wattle**



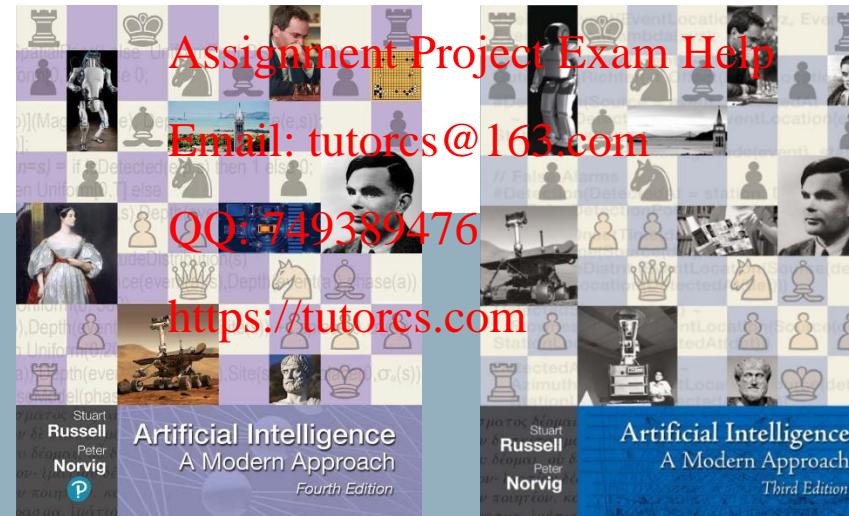
Australian
National
University

程序代写代做 CS编程辅导

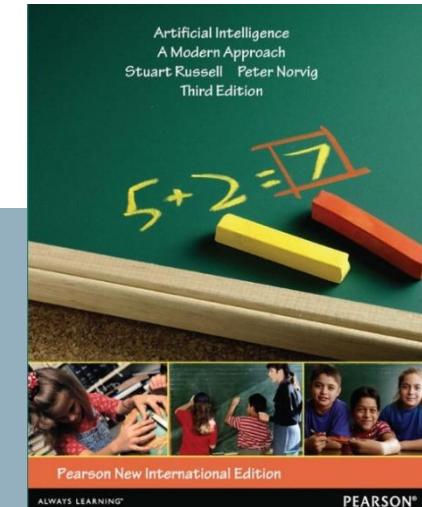


Intelligent Agents

WeChat: cstutorcs



Chapter 2



PEAS

程序代写代做 CS 编程辅导

To design a rational agent, you must specify the **task environment**



Consider, e.g., the task of ~~driving~~ **hailing a driverless taxi**:

- **Performance measure:** WeChat: cstutorcs
– safety, destination, profits, legality, comfort, ...
- **Environment:** Assignment Project Exam Help
– streets/freeways, traffic, pedestrians, weather, ...
- **Actuators:** QQ: 749389476
– steering, accelerator, brake/horn, blinkers, ...
- **Sensors:** https://tutorcs.com
– GPS, video, accelerometers, gauges, engine sensors, ...

Properties of Task Environments

程序代写代做 CS 编程辅导

- **Fully vs partially observable**
 - do the agent sensors give access to relevant information about the environment state?
- **Deterministic vs stochastic**
 - is the next state completely determined by the current state and executed action?
- **Known vs unknown**
 - does the agent know the environment's laws of physics?
- **Episodic vs sequential**
 - is the next decision independent of the previous ones?
- **Static vs dynamic**
 - can the environment change whilst the agent is deliberating?
 - **Semi-dynamic:** only the performance score changes.
- **Discrete vs continuous**
 - can time, states, actions, percepts be represented in a discrete way?
- **Single vs multi-agent**
 - is a single agent making decisions, or do multiple agents need to compete or cooperate to maximise interdependent performance measures?



WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Agent types

程序代写代做CS编程辅导



- Four basic types of agents in a series of increasing generality:
 - simple reflex agents
 - reflex agents with state
 - goal-based agents
 - utility-based agents
- All these can be turned into learning agents

WeChat: cstutorcs

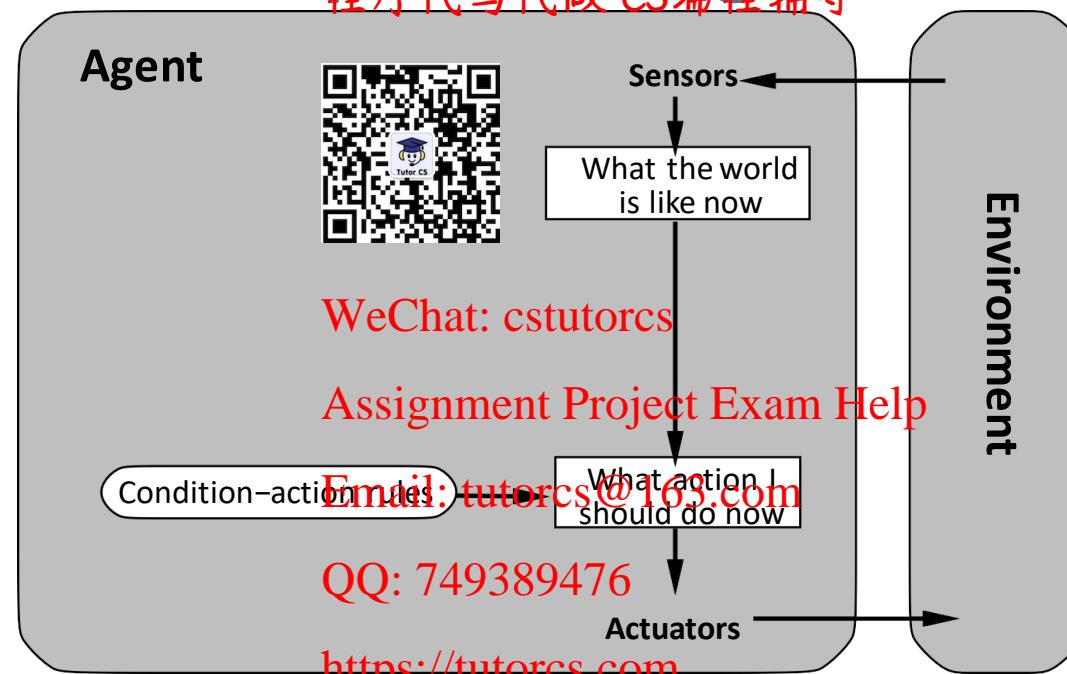
Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Simple reflex agents



Decisions are made based on the **current percept** only. Raises issues for partially observable environments.

Example

程序代写代做CS编程辅导

```
function REFLEX-VACUUM-CONTROLLER(location,status) returns an action
  if status = Dirty then return Suck
  else if location = A then return Right
  else if location = B then return Left
```



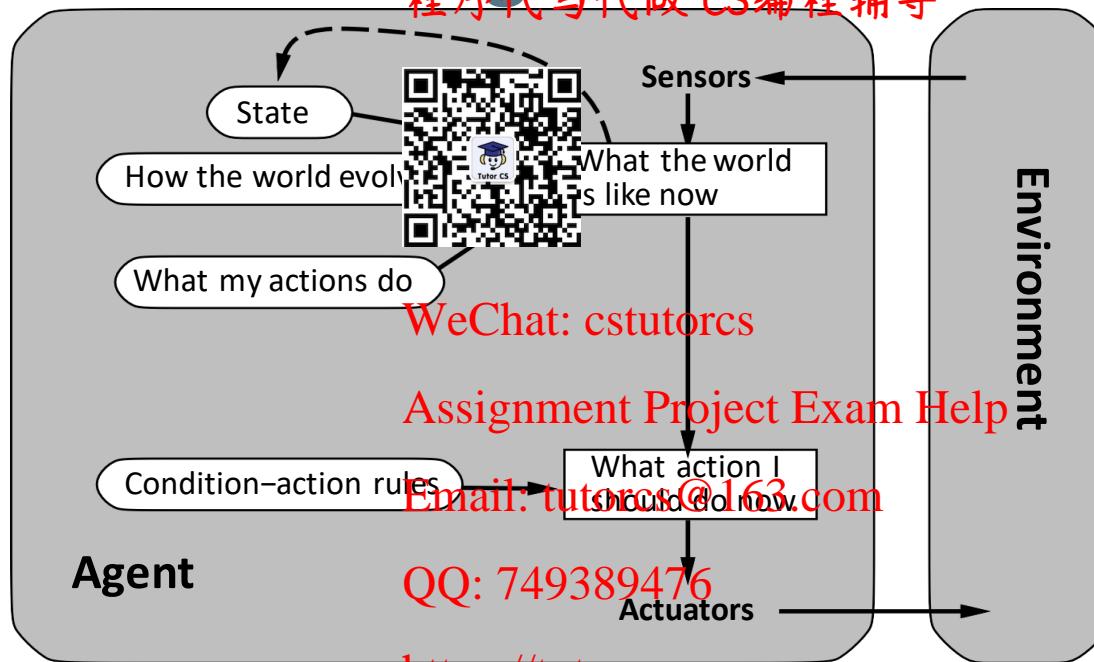
Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Reflex agents with state



The **internal state** keeps track of relevant unobservable aspects of the environment. The **environment model** describes how the environment works (how the environment state is affected by actions)

Example

程序代写代做CS编程辅导

```
function VACUUM-AGENT-WITH-STATE( (location, status) ) returns an action
static: last_A, last_B, n = 0, time = 0, tmax = infinity
increment last_A and
if location = A then
else last_B = 0
case
status = Dirty: WeChat: cstutorcs
    return Suck
location = A: Assignment Project Exam Help
    if last_B > 3 then return Right
    else return NoOp
location = B:
    if last_A > 3 then return Left
    else return NoOp
```

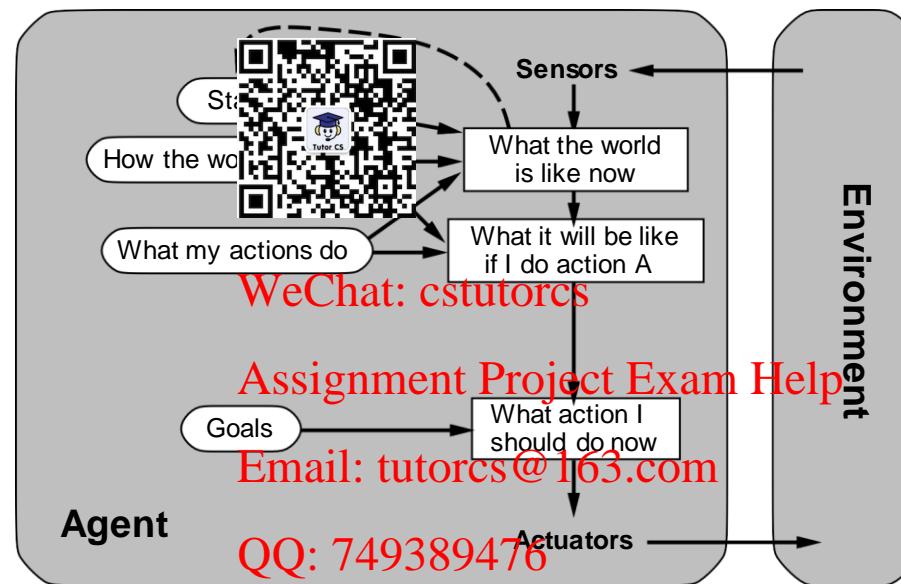


<https://tutorcs.com>

The time passed since a location was visited is a proxy for the likelihood of this location's status changing from clean to dirty.

Goal-Based agents

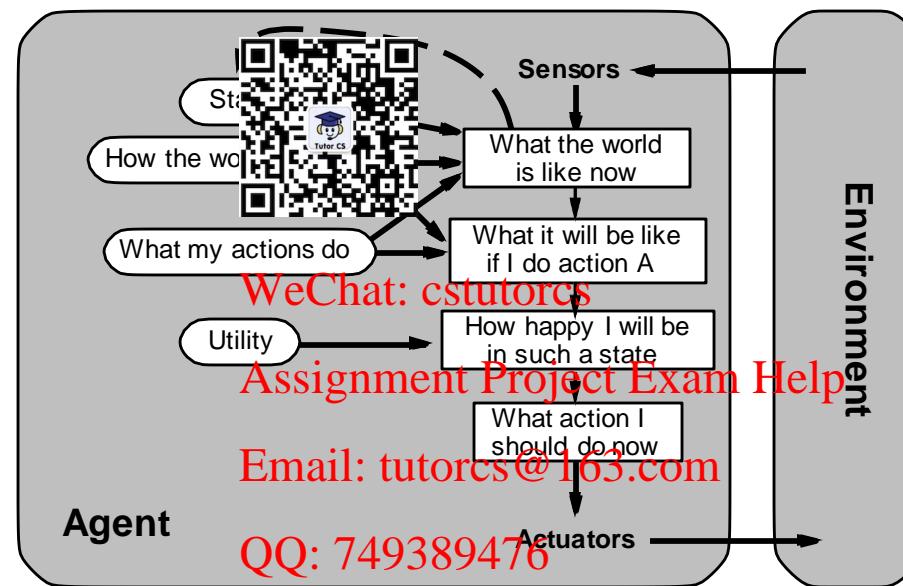
程序代写代做 CS编程辅导



- The **goal** describes desirable situations. <https://tutorcs.com>
- The agent combines goal and environment model to choose actions.
- **Planning** and **search** are AI subfields devoted to building goal-based agents.

Utility-based agents

程序代写代做 CS编程辅导



- The **utility function** internalises the performance measure.
- Under uncertainty, the agent chooses actions that maximise the expected utility.

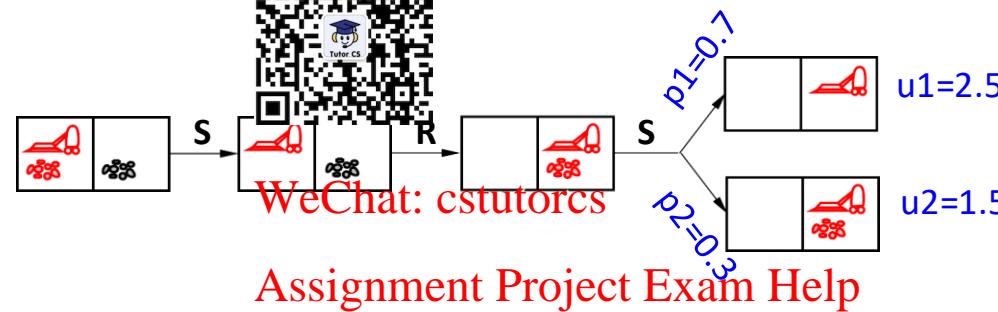
Utility-based agents

程序代写代做 CS编程辅导

Rational agent: chooses the action that maximises expected utility:

Utility function:

1 per clean room
-0.5 per action



Expected utility of *Suck*:

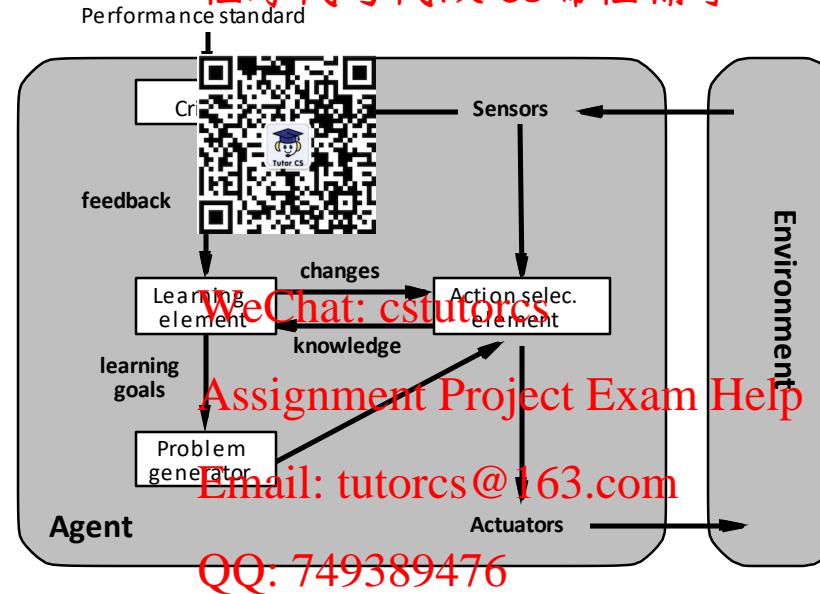
$$p_1 \times u_1 + p_2 \times u_2 = 0.7 \times 2.5 + 0.3 \times 1.5 = 2.2$$

QQ: 749389476

- *Suck* has an expected utility of 2.2 <https://tutorcs.com>
- *NoOp* has an expected utility of 2
- *Left* has an expected utility of 1.5

Learning agents

程序代写代做 CS 编程辅导



- The **action selection** element is what we described earlier.
- The **learning** element uses feedback from the **critic** to modify the action selection.
- The **problem generator** suggests actions that lead to new informative experience.

Exploration vs Exploitation

程序代写代做 CS 编程辅导

A fundamental dilemma for learning agents:

- **Exploitation**: greedily uses what the agent has learnt to select the action that will, in the light of the current knowledge, have the best outcome
- **Exploration**: taking some other (possibly random) action to learn more, hoping to find something even better than what is currently known

In practice, agents must explore to avoid getting stuck in severely sub-optimal behaviour, but exploration has a cost.

Typically, a smart agent explores more in early stages than later on



WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>



Australian
National
University

程序代写代做 CS编程辅导



Goal-based Agents: Solving Problems by Searching

WeChat: cstutorcs

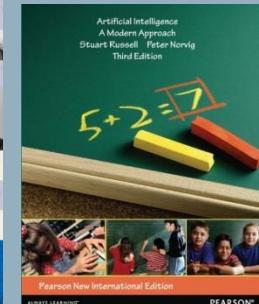
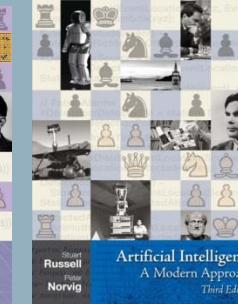
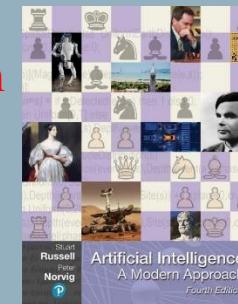
Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

Chapter 3, Sections 1-4

<https://tutorcs.com>



Search for Atari Games

程序代写代做CS编程辅导



WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Lipovetzky, Nir, Miquel Ramirez, and Hector Geffner. "Classical planning with simulators: Results on the atari video games." In 24th International Joint Conference on Artificial Intelligence. 2015.

<https://www.youtube.com/watch?v=P-603qPMkSg>

Problem solving agents

程序代写代做 CS编程辅导

- Form of **goal-based agent** that formulates the problem of **reaching a goal** in its environment, searches for a **sequence of actions** solving the problem, and executes it.
- Assumptions about the task environment:
 - fully observable
 - deterministic
 - static
 - single agent
- **Offline** (or open-loop) problem solving is suitable under those assumptions; the entire sequence solution can be **executed “eyes closed.”**



WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Outline

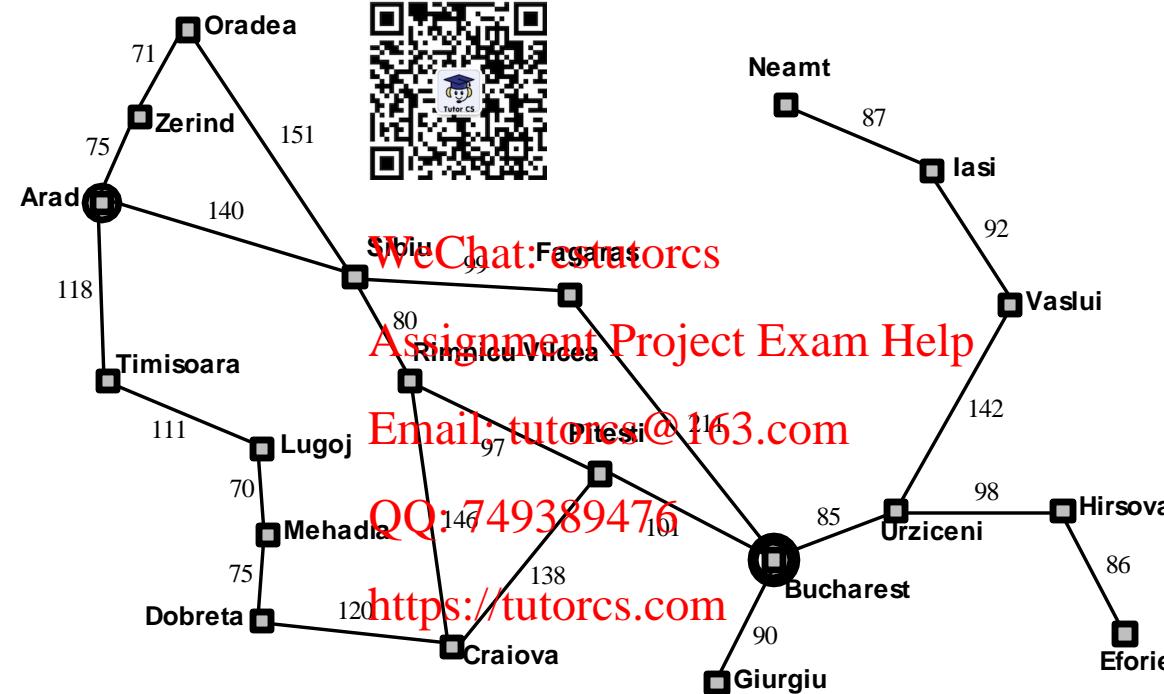
程序代写代做 CS 编程辅导



- Problem formulation
- Problem formulation examples WeChat: cstutorcs
- Tree search algorithm Assignment Project Exam Help
- Uninformed search strategies Email: tutorcs@163.com
- Informed search strategies QQ: 749389476
<https://tutorcs.com>

Example: Romania

程序代写代做 CS编程辅导



Example: Romania

程序代写代做 CS编程辅导

- On holiday in Romania
 - currently in Arad
 - flight leaves tomorrow from Bucharest



- **Formulate problem:**
 - **initial state:** in Arad
 - **goal:** be in Bucharest
 - **states:** various cities
 - **actions:** drive between cities

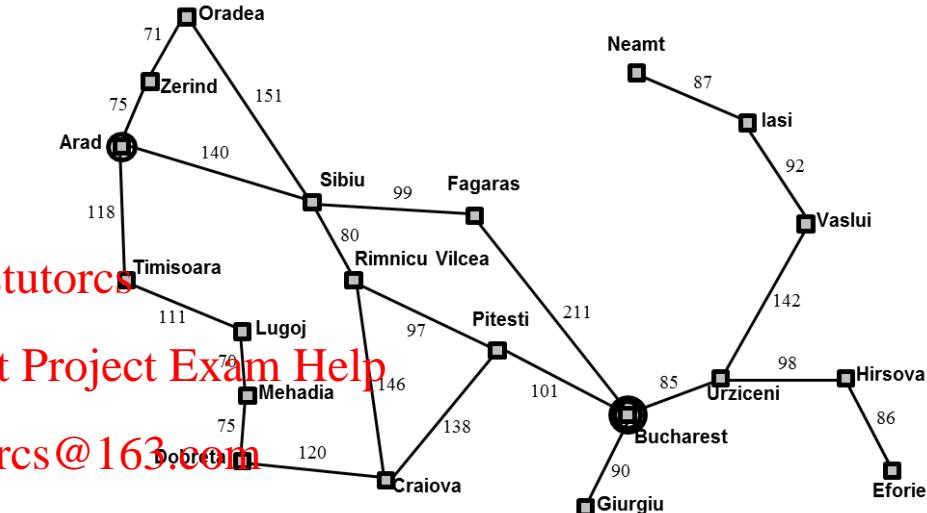
WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

- **Search for a solution:** <https://tutorcs.com>
 - sequence of drive actions or equivalently (in this case) sequence of cities, e.g. Arad, Sibiu, Fagaras, Bucharest



Problem formulation

程序代写代做CS编程辅导

A problem is defined by four items:

- initial state: e.g., “at Arad”
- successor function: $S(x) = \{ \langle x, y \rangle \mid \text{transition from } x \text{ to } y \}$ n-state pairs
e.g., $S(Arad) = \{ \langle Arad \rightarrow Zerind \rangle, \langle Arad \rightarrow Sibiu, Sibiu \rangle, \langle Arad \rightarrow Timisoara, Timisoara \rangle \}$

WeChat: cstutorcs

- goal test, can be

- explicit, e.g., if $x = \text{"at Bucharest"}$
- implicit, e.g., if $\text{HasAirport}(x)$ is True

Assignment Project Exam Help

Email: tutorcs@163.com

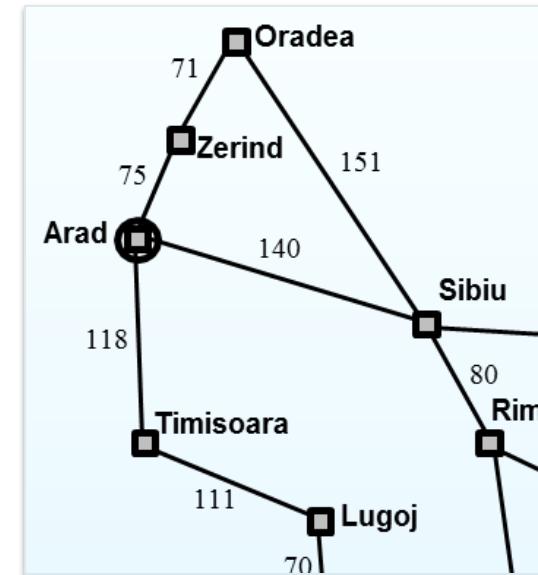
- path cost (additive)

QQ: 749389476

- e.g., sum of distances, number of actions executed, etc.

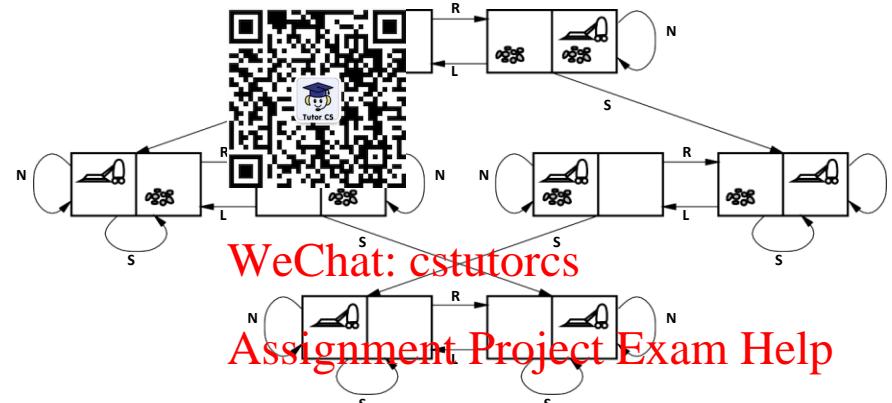
- $c(x, a, y) \geq 0$ is the step cost of going from state x to state y via action a .

- A solution is a sequence of actions leading from the initial state to a goal state



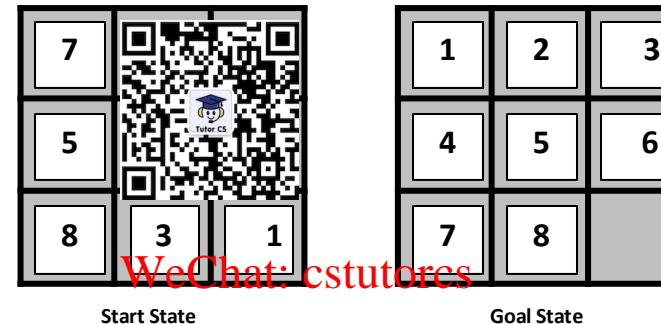
Example: Vacuum world

程序代写代做 CS 编程辅导



- states? dirt presence in each room and robot location (ignore dirt amounts)
- actions? *Left, Right, Suck, NoOp*
- goal test? no dirt
- path cost? 1 per action (0 for *NoOp*)

Example: The 8-puzzle



- **states**? Integer locations of tiles (ignore intermediate positions)
 - **actions**? move blank left, right, up, down (ignore unjamming etc.)
 - **goal test**? goal state (given)
 - **path cost**? 1 per move

[Note: optimal solution of n -Puzzle family is NP-hard]

Selecting a state space

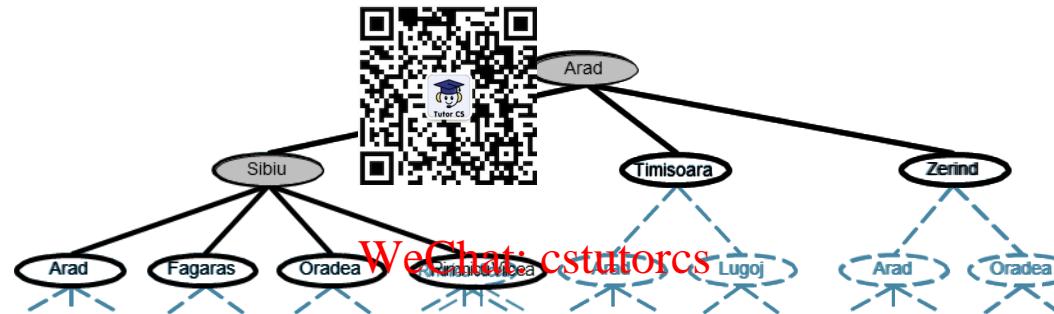
程序代写代做 CS编程辅导

- Real world is absurdly complex
⇒ state space must be abstract for problem solving
- (Abstract) state = set of real states
- (Abstract) action = complex combination of real actions
 - e.g., “Arad → Zerind” represents a complex set of possible routes, detours, rest stops, etc.
 - For guaranteed realizability, any real state “in Arad” must get to some real state in “in Zerind”.

Assignment Project Exam Help
WeChat: cstutors
Email: tutorcs@163.com
QQ: 749389476
- (Abstract) solution = set of real paths that are solutions in the real world
<https://tutorcs.com>
- Abstraction should be “easier” than the original problem!

Tree search example

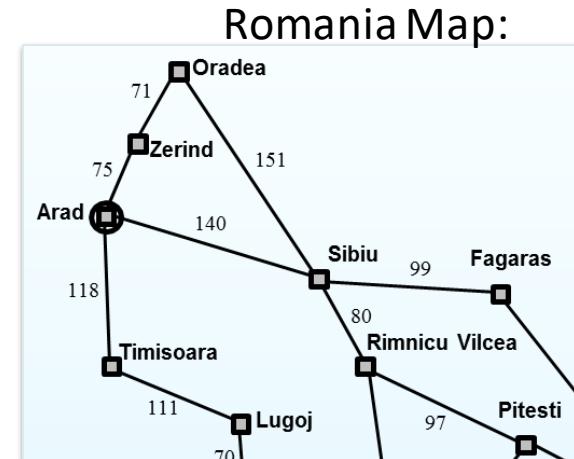
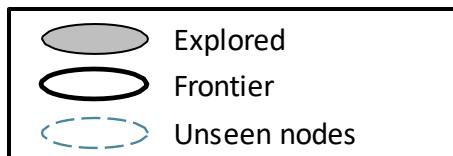
程序代写代做 CS编程辅导



Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>



Tree search algorithm

程序代写代做 CS编程辅导

Basic idea:

- offline, simulated exploration
• generate space by generating successors of already-explored nodes (a.k.a. **expanding nodes**)



WeChat: cstutorcs

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
    initialize the search tree using the initial state of problem
    Assignment Project Exam Help
    loop do
        if there are no candidates for expansion on the frontier then return failure
        choose a frontier node for expansion according to strategy
        Email: tutorcs@163.com
        QQ: 749389476
        if the node contains a goal state then return the corresponding solution
        else expand the node and add the resulting nodes to the frontier of the tree
    end
```

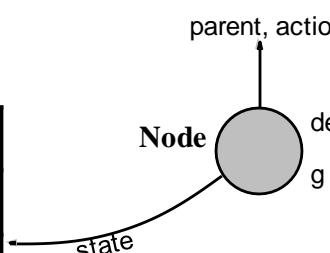
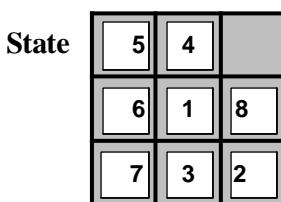
Implementation: states vs. nodes

程序代写代做 CS 编程辅导

- A **state** is a representation of a physical configuration
- A **node** is a data structure containing part of a search tree
 - includes: state, parent, children, depth, **path cost** $g(n)$
- States **do not have** parents, children, depth, or path cost!

WeChat: cstutorcs

8-puzzle:

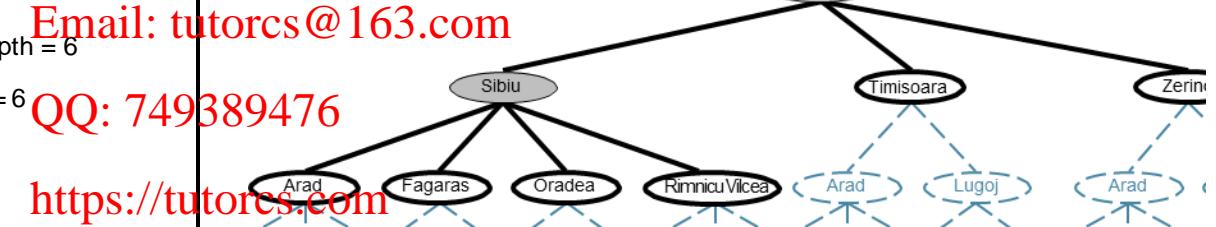


Romania: Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>



Implementation: Expand and Frontier

程序代写代做 CS 编程辅导

- The **EXPAND** function creates new nodes, filling in the various fields and using the **SUCCESSORFN** of the problem to get the corresponding states.
 - Example: **SUCCESSORFN(Arad)** → $\{ \langle Arad \rightarrow Zerind, Zerind \rangle, \langle Arad \rightarrow Sibiu, Sibiu \rangle, \langle Arad \rightarrow Timisoara, Timisoara \rangle \}$



- Frontier implemented as **priority queue** of nodes ordered according to strategy

Implementation: general tree search

程序代写代做CS编程辅导



WeChat: cstutorcs

function EXPAND(*node*, *problem*) returns a set of nodes
 successors \leftarrow the empty set
 for each *action*, *result* \in SUCCESSOR-F(*problem*, STATE[*node*]) **do**
 s \leftarrow a new NODE
 PARENT-NODE[*s*] \leftarrow *node*; ACTION[*s*] \leftarrow *action*; STATE[*s*] \leftarrow *result*
 PATH-COST[*s*] \leftarrow PATH-COST[*node*] + STEP-COST(STATE[*node*], *action*,
 result)
 DEPTH[*s*] \leftarrow DEPTH[*node*] + 1
 add *s* to *successors*
 return *successors*

Assignment Project Exam Help
Email: tutors@163.com
QQ: 749389476
<https://tutorcs.com>



Australian
National
University

程序代写代做 CS编程辅导



Uniform Search Strategies

WeChat: cstutorcs

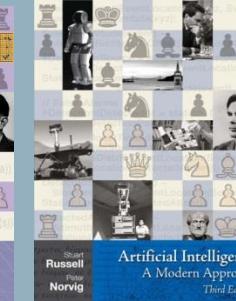
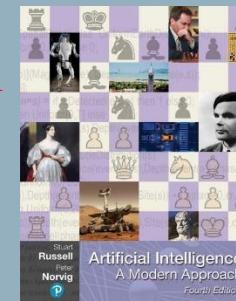
Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

Chapter 3, Section 4

<https://tutorcs.com>



Uninformed search strategies

程序代写代做 CS 编程辅导

- A strategy is defined by picking the order of node expansion.
- This is the order used for the **frontier queue** implementing the frontier.
- **Uninformed** strategies use **only** the information available in the definition of the problem:
 - Breadth-first search
 - Uniform-cost search
 - Depth-first search
 - Depth-limited search
 - Iterative deepening search

WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>



Breadth-first search

程序代写代做 CS编程辅导

- **Strategy:** expand shallowest unexpanded node

- **Implementation:**

- **frontier** is a FIFO queue, i.e., new successors go at end

Depth

WeChat: cstutorcs

0

Assignment Project Exam Help

1

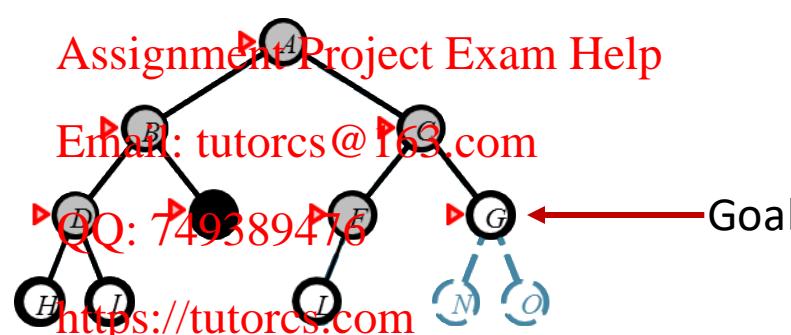
Email: tutorcs@163.com

2

QQ: 749389476

3

<https://tutorcs.com>

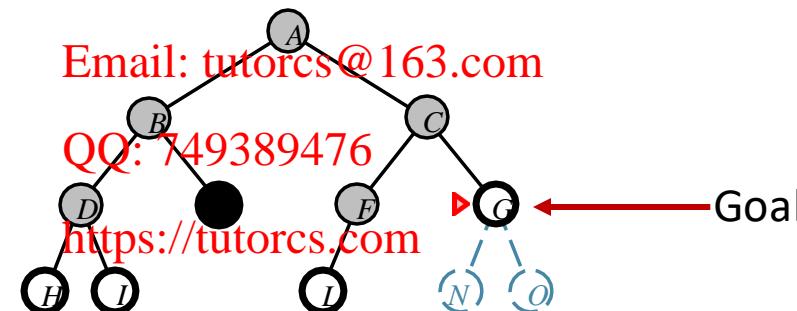


Evaluating Search Strategies

程序代写代做 CS 编程辅导

- Strategies are evaluated along the following dimensions:
 - completeness - does it always find a solution if one exists?
 - solution optimality - does it always find a least-cost solution?
 - time complexity - number of nodes generated (or expanded)
 - space complexity - maximum number of nodes in memory

Assignment Project Exam Help

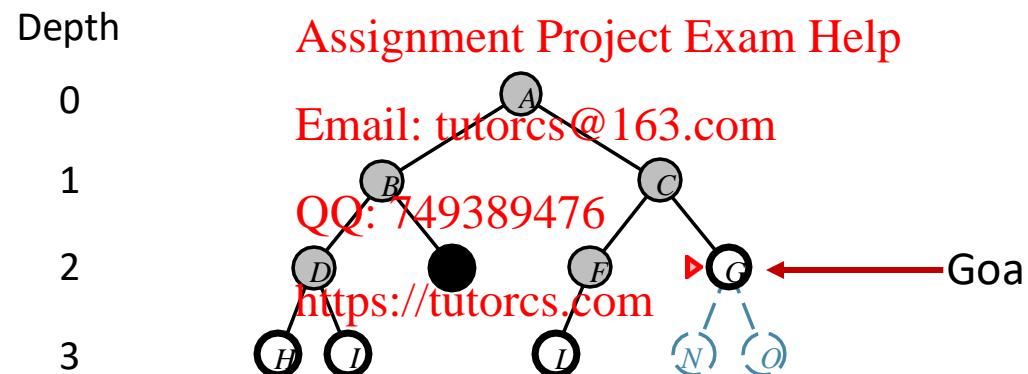


Time and Space Complexity

程序代写代做 CS 编程辅导

- Time and space complexity are measured using big-O notation in terms of
 - b : maximum **branching factor** of the search tree
 - d : **depth** of the shallowest goal node
 - m : **maximum** depth of the state space (may be ∞)

WeChat: cstutorcs

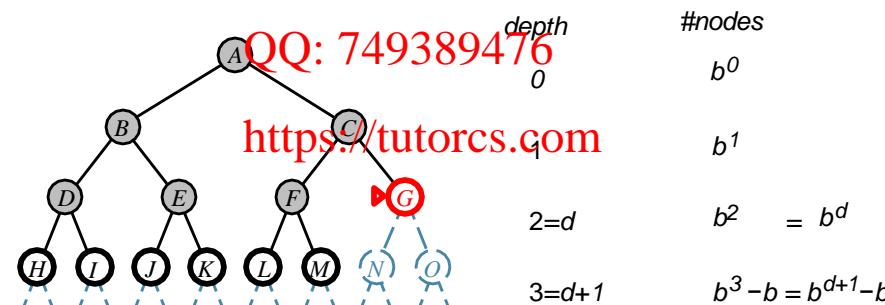


Properties of breadth-first search

程序代写代做 CS 编程辅导

- Complete? Yes (if b and d are finite)
- Time? $1 + b + b^2 + b^3 + \dots + b^{d+1} - b = O(b^{d+1})$, i.e., exp. in d
- Space? $O(b^{d+1})$
 - if b and d are finite $b = 10$, 1 million node/sec, 1Kb/node, $d = 12$ would take 13 days and 1 petabyte of memory.
- Optimal? Yes if cost = 1 per step, not optimal in general

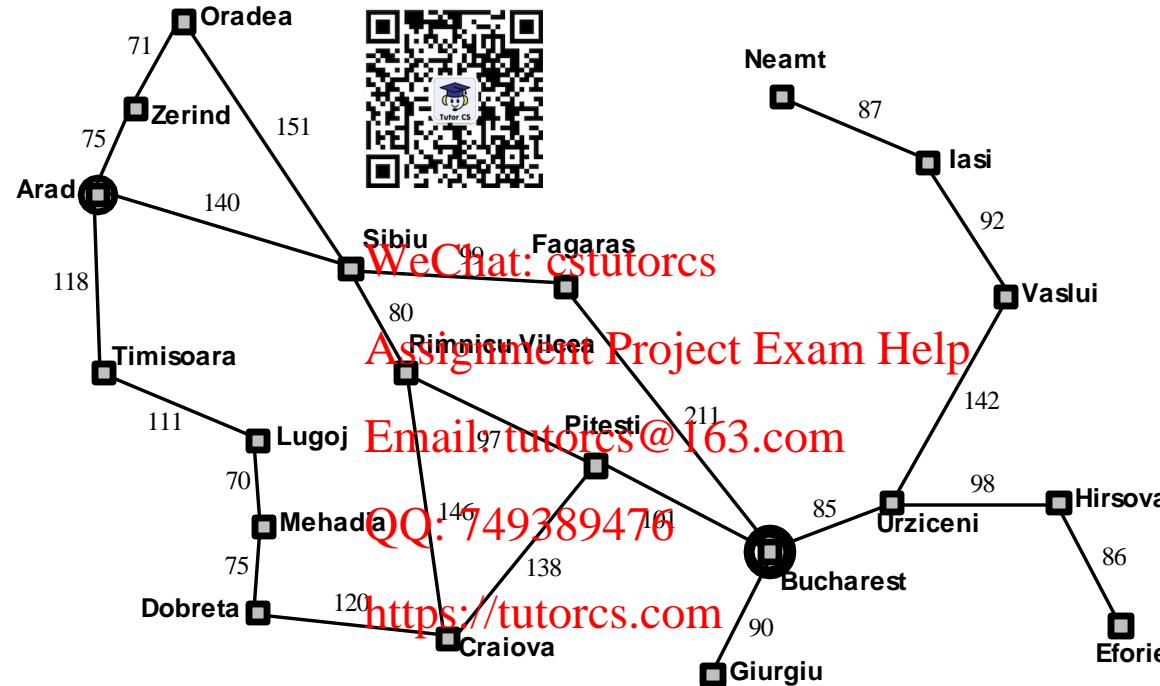
WeChat: cstutorcs
Assignment Project Exam Help
Email: tutorcs@163.com



b : max. branching factor
 d : depth of the shallowest solution
 m : max. depth of the state space

Example: Romania

程序代写代做 CS编程辅导



Uniform-cost search

程序代写代做CS编程辅导

- **Strategy:** Expand least-cost unexpanded node
- **Implementation:**
 - *frontier* = queue ordered by cost (i.e., $g(n)$), lowest first
- Equivalent to breadth-first if step costs all equal
WeChat: cstutorcs
- Complete? Yes, if step cost $\Delta \leq \epsilon > 0$
Assignment Project Exam Help
- Time? # of nodes with $g \leq C^*$, $O(b^{1+[C^*/\epsilon]})$
Email: tutorcs@163.com
where C^* is the cost of the optimal solution
QQ: 749389476
- Space? # of nodes with $g \leq C^*$, $O(b^{1+[C^*/\epsilon]})$
<https://tutorcs.com>
- Optimal? Yes—nodes expanded in increasing order of $g(n)$

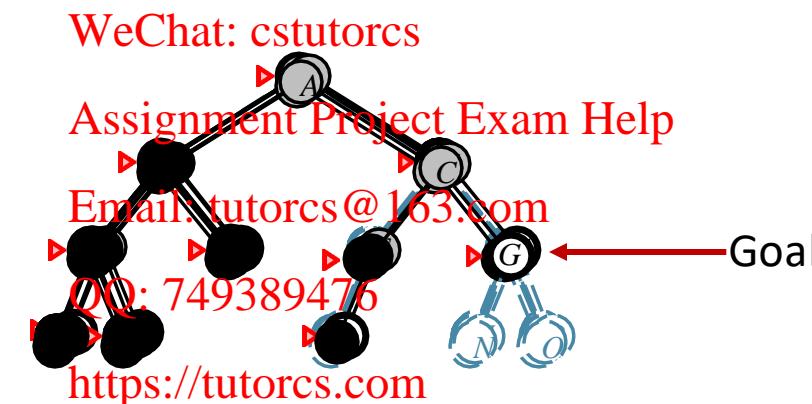


b : max. branching factor
 d : depth of the shallowest solution
 m : max. depth of the state space

Depth-first search

程序代写代做 CS 编程辅导

- **Strategy:** Expand deepest unexpanded node
- **Implementation:**
 - *frontier* = LIFO queue, i.e., put $s \in S$ at front

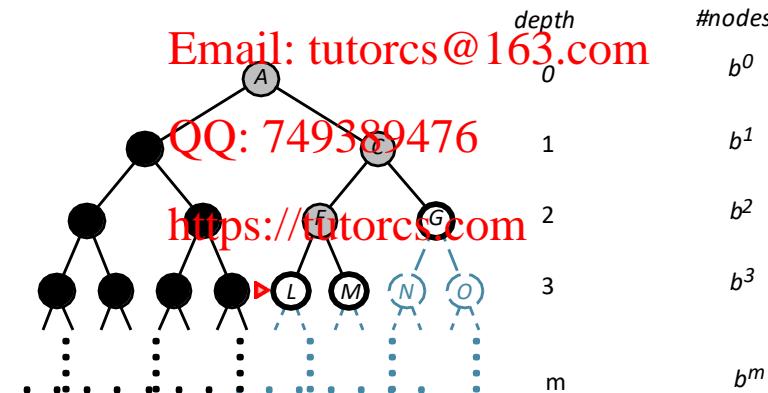


Properties of depth-first search

程序代写代做 CS 编程辅导

- Complete? No: fails in infinite-depth spaces, spaces with loops
 ⇒ complete in finite-depth spaces if we modify to avoid repeated states along path
- Time? $O(b^m)$: terrible if m is much larger than d
 ⇒ if solutions are dense, We may be much faster than breadth-first
- Space? $O(bm)$, i.e., linear space! (deepest node + ancestors + their siblings)
- Optimal? No

Email: tutorcs@163.com



QQ: 749389476
<https://tutorcs.com>

b : max. branching factor
 d : depth of the shallowest solution
 m : max. depth of the state space

Breadth-first versus depth-first search

程序代写代做 CS 编程辅导

- Use breadth-first search when there exists **short** solutions.
- Use depth-first search when there exists **many** solutions.



Eternity II Puzzle

2 million dollar prize! Few deep solutions (presumably...)

Iterative deepening search

程序代写代做 CS 编程辅导

- Combines advantages of breadth-first and depth-first search:
 - is complete
 - returns shallowest solution
 - uses linear amount of memory



WeChat: cstutorcs

- Performs a series of depth limited depth-first searches

Assignment Project Exam Help

```
function ITERATIVE-DEEPENING-SEARCH(problem) returns solution/failure
    Email: tutorcs@163.com
    inputs: problem, a problem
    QQ: 749389476
    for depth  $\leftarrow$  0 to  $\infty$ 
        result  $\leftarrow$  DEPTH-LIMITED-SEARCH(problem, depth)
        if result  $\neq$  cutoff then return result
    end
```

<https://tutorcs.com>

Depth-limited search

程序代写代做 CS 编程辅导

- Depth-first search with depth limit l , i.e., nodes at depth l have no successors
- **Recursive implementation**



```
function DEPTH-LIMITED-SLS(problem, limit) returns soln/fail/cutoff
    RECURSIVE-DLS(MAKE-NODE(INITIAL-STATE[problem]), problem, limit)
```

```
function RECURSIVE-DLS(node, problem, limit) returns soln/fail/cutoff
    cutoff-occurred? ← false
    if GOAL-TEST(problem, STATE[node]) then return node
    else if DEPTH[node] > limit then return cutoff
    else for each successor in EXPAND(node, problem) do
        result ← RECURSIVE-DLS(successor, problem, limit)
        if result = cutoff then cutoff-occurred? ← true
        else if result ≠ failure then return result
    if cutoff-occurred? then return cutoff else return failure
```

<https://tutorcs.com>

- **solution**: solution found within depth limit
- **cutoff**: no solution within the depth limit
- **failure**: the problem has no solution

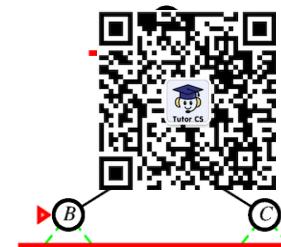
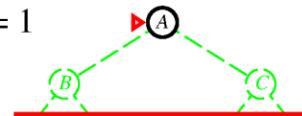
Iterative deepening search – Example

程序代写代做 CS编程辅导

Limit = 0

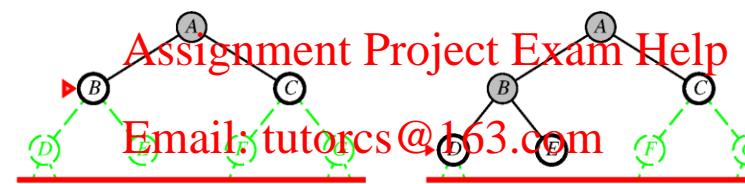
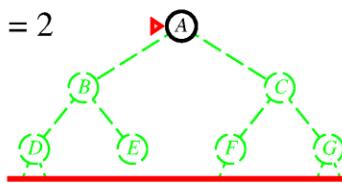


Limit = 1



WeChat: cstutorcs

Limit = 2

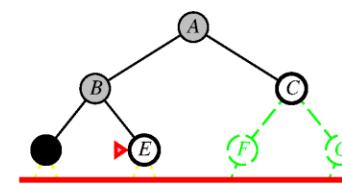
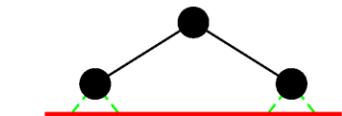
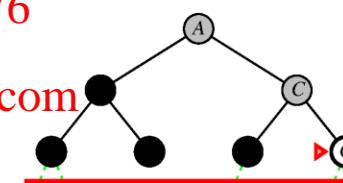
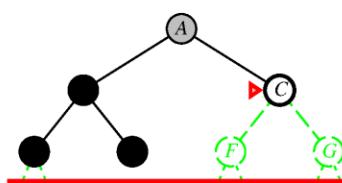


Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>



Properties of iterative deepening search

程序代写代做CS编程辅导

- Complete? Yes (if b and d are finite)
- Time? $(d + 1)b^0 + db^1 + (d + 1)b^2 + \dots + b^d = O(b^d)$
- Space? $O(bd)$
- Optimal? Yes, if step cost = 1 (can be modified to explore uniform-cost tree)



b : max. branching factor
 d : depth of the shallowest solution
 m : max. depth of the state space

WeChat: cstutorcs
 Assignment Project Exam Help

Numerical comparison for $b = 10$ and $d = 5$, solution at far-right leaf:

Email: tutorcs@163.com

Time: IDS doesn't do much worse than BFS

$$N(BFS) = 1 + 10 + 100 + 1,000 + 10,000 + 100,000 = 111,111$$

$$N(IDS) = 6 + 50 + 400 + 3,000 + 20,000 + 100,000 = 123,456$$

<https://tutorcs.com>

Assumes BFS applies the goal test when a node is generated

Space: IDS does much better $N(IDS) = 50$, $N(BFS) \approx 1,000,000$

Summary of algorithms

程序代写代做CS编程辅导

Criterion	Breadth-First		Depth-First	Depth-Limited	Iterative Deepening
Complete?	Yes ¹		No	Yes, if $l \geq d$	Yes ¹
Time	b^{d+1}	$b^{1+[C^*/\epsilon]}$	b^m	b^l	b^d
Space	b^{d+1}	$b^{1+[C^*/\epsilon]}$	bm	bl	bd
Optimal?	Yes ²	Yes	No	No	Yes ²

1. Only if b and d are finite.

QQ: 749389476

2. Only if all step costs are identical

<https://tutorcs.com>

b : max. branching factor
 d : depth of the shallowest solution
 m : max. depth of the state space



Australian
National
University

程序代写代做 CS编程辅导



Information search algorithms

WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

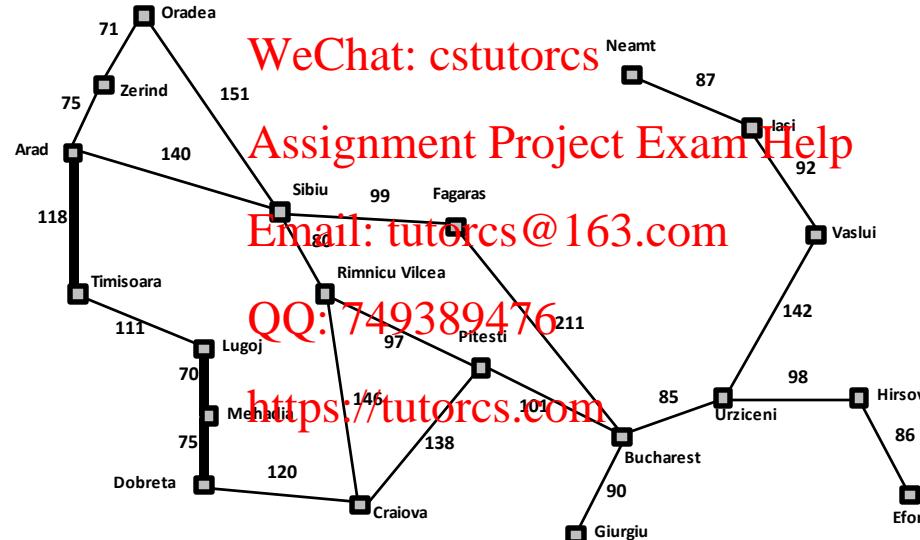
Chapter 3, Sections 5–6

<https://tutorcs.com>

Heuristic Function

程序代写代做 CS 编程辅导

- Estimates the cost from a given state to the goal
 - Estimate for the Travel in Romania?
 - Straight Line Distance



Straight-line
dist. to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	178
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	100
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

Outline

程序代写代做 CS 编程辅导



- Evaluation functions (use heuristics)
- Greedy search
- A* search

WeChat: cstutorcs

Assignment Project Exam Help

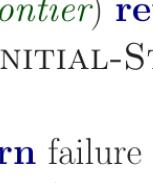
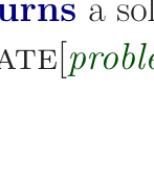
Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Review: Tree search

程序代写代做 CS编程辅导

```
function TREE-SEARCH(, frontier) returns a solution, or failure
    frontier ← INSERT(, INITIAL-STATE[problem]), frontier)
    loop do
        if frontier is empty  then return failure
        node ← REMOVE-FRONT(frontier)
        if GOAL-TEST(problem, STATE(node)) then return node
        frontier ← INSERTALL(EXPAND(node, problem), frontier)



```

WeChat: estutorcs
Assignment Project Exam Help

- The goal test is performed when the node is popped from the frontier,
NOT when it is generated during expansion.
– This is important when looking for optimal solutions.
<https://tutorcs.com>
- A strategy is defined by picking the order of node expansion

Evaluation function

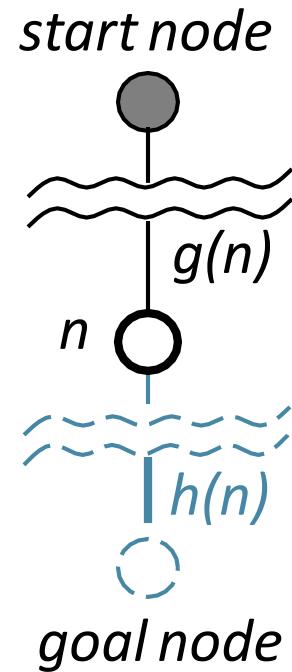
程序代写代做 CS 编程辅导

- Evaluation function $f(n) = g(n) + h(n)$
 - estimate of “desirability”, user-defined, problem-specific.
 - $g(n)$ = cost so far to reach n
 - $h(n)$ = estimated cost from n to the closest goal (heuristic)
 - $f(n)$ = estimated total cost of path through n to goal
 - The lower $f(n)$, the more desirable n is
- Implementation:
 - $frontier$ is a queue sorted in ascending value of $f(n)$
- Special cases:
 - uniform cost search (uninformed): <https://tutorcs.com>
 - greedy search (informed): $f(n) = h(n)$
 - A* search (informed): $f(n) = g(n) + h(n)$

WeChat: cstutorcs
Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476



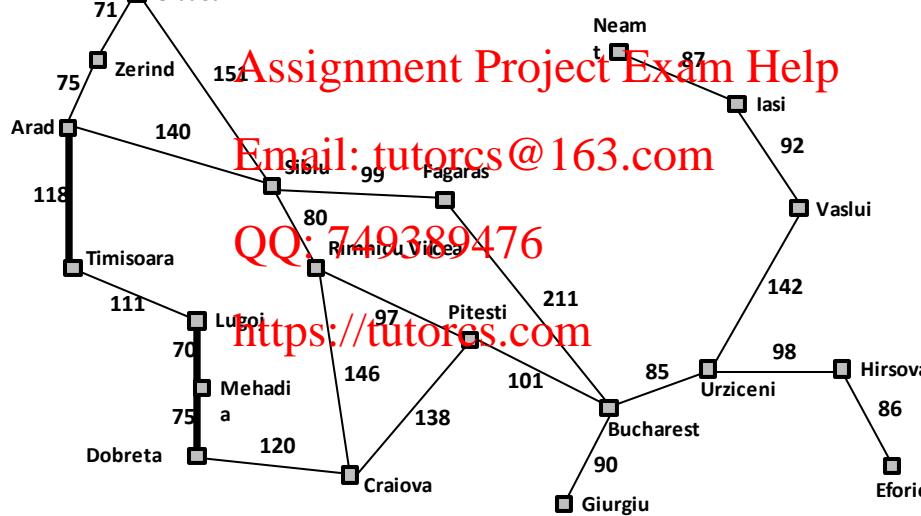
Greedy search

程序代写代做 CS编程辅导

- Evaluation function $f(n) = h(n)$ (entirely heuristic)
= estimate of cost from n to the closest goal
- E.g., $h_{SLD}(n)$ = straight-line distance from n to Bucharest
- Greedy search expands the node that **appears** to be closest to goal



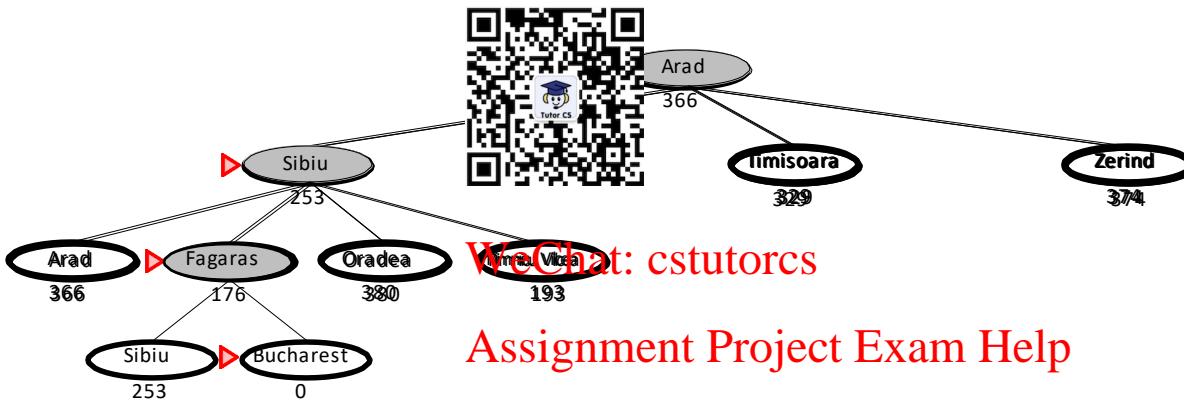
cstutorcs



Straight-line dist to Bucharest	
Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	178
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	100
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

Greedy Search Example

程序代写代做 CS编程辅导



Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

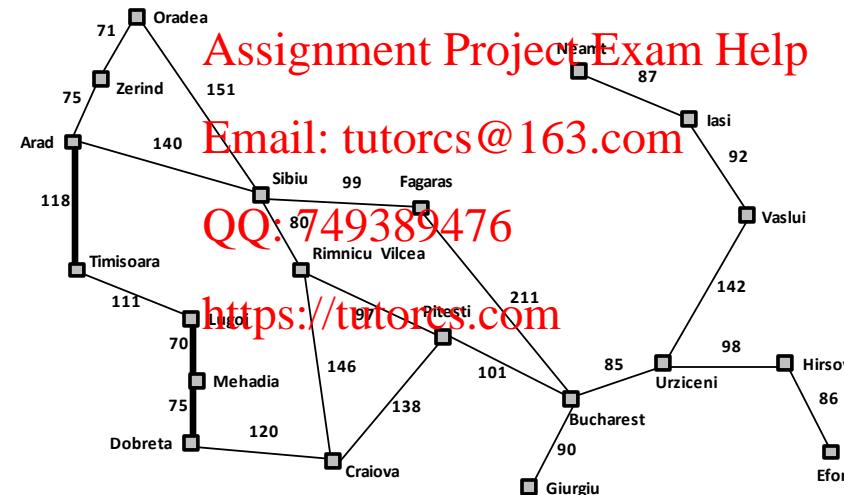
Straight-line
dist. to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	178
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	100
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

Properties of greedy search

程序代写代做 CS编程辅导

- Complete? No. It can get stuck in loops, e.g., with going from lasi to Fagaras:
 lasi → Neamt → lasi → Neamt → ...
- Complete in finite space  via repeated-state checking
- Time? $O(b^m)$, but a good heuristic can give dramatic improvement
- Space? $O(b^m)$
- Optimal? No



b : max. branching factor
 d : depth of the shallowest solution
 m : max. depth of the state space

A* Search

程序代写代做CS编程辅导

- Idea: avoid expanding paths that are already expensive
- Evaluation function $f(n) = g(n) + h(n)$
 - $g(n)$ = cost so far to reach n
 - $h(n)$ = estimated cost from n to the closest goal
 - $f(n)$ = estimated total cost of path through n to goal
- Admissible heuristic:
 - $\forall n \ h(n) \leq h^*(n)$ where $h^*(n)$ is the true cost from n
 - Also require $h(n) \geq 0$, so $h(G) = 0$ for any goal G
- E.g., $h_{SLD}(n)$ never overestimates the actual road distance
- When $h(n)$ is admissible, $f(n)$ never overestimates the total cost of the optimal path through n to the goal
- Theorem: if h is admissible, A* search finds the optimal solution

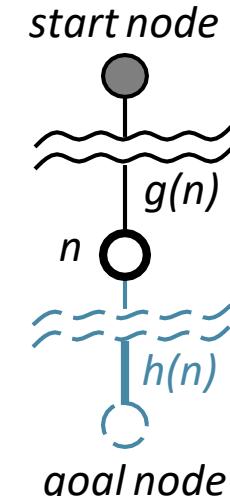


WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

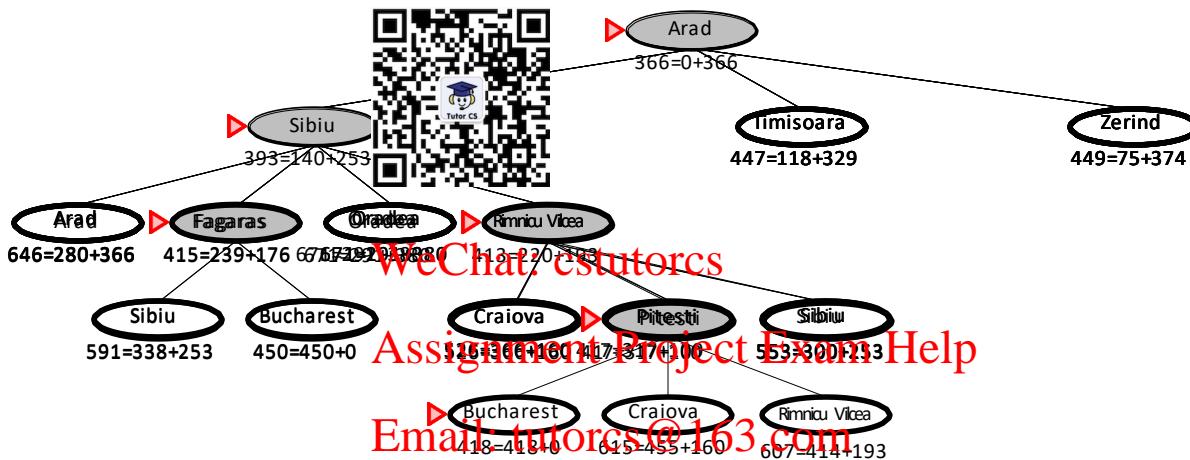
QQ: 749389476



A* Search Example

程序代写代做 CS编程辅导

$$f(n) = g(n) + h(n)$$



WeChat: csutorcs
 Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

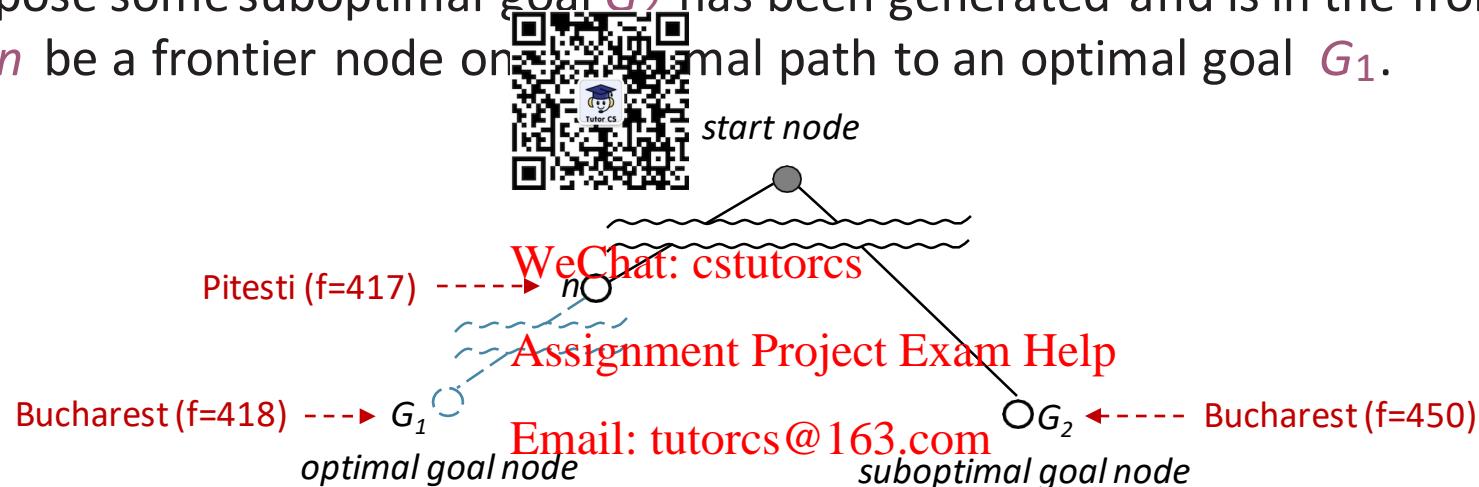
Straight-line
dist. to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	178
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	100
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

Optimality of A* (based on admissibility)

程序代写代做 CS 编程辅导

- Suppose some suboptimal goal G_2 has been generated and is in the frontier. Let n be a frontier node on the optimal path to an optimal goal G_1 .



$$f(G_2) = g(G_2) \text{ since } G_2 \text{ is a goal node, hence, } h(G_2) = 0$$

$$> g(G_1) \text{ since } G_1 \text{ is suboptimal}$$

$$\geq f(n) \text{ since } h \text{ is admissible: } f(n) = g(n) + h(n) \leq g(n) + h^*(n) = g(G_1)$$

- Since $f(G_2) > f(n)$, A* will never select G_2 for expansion