

COMP3620 / 6320 – S1 2023

程序代写代做 CS编程辅导



This lecture will be recorded

Turn off your camera if you do not want to be in the recording

WeChat: cstutorcs

If you cannot see my whole slide, go to **View Options -> Zoom Ratio** and select **Fit to Window**

Assignment Project Exam Help

Email: tutorcs@163.com

If you want to ask a question, either:

- Raise your hand and I will open your mic, or
- Type the question in the chat box

QQ: 749389476

<https://tutorcs.com>

Outline for Today

程序代写代做 CS编程辅导



- Finish A*
- Generalizing Tree Search to Graph Search
- How to construct (Admissible) Heuristics
- Adversarial Search

WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Recap: A* Search

程序代写代做 CS编程辅导

- **Idea**: avoid expanding paths that are already expensive
- Evaluation function $f(n) = g(n) + h(n)$
 - $g(n)$ = cost so far to reach n
 - $h(n)$ = estimated cost from n to the closest goal
 - $f(n)$ = estimated total cost of path through n to goal

WeChat: cstutorcs

- **Admissible** heuristic:

- $\forall n, h(n) \leq h^*(n)$ where $h^*(n)$ is the **true** cost from n
- Also require $h(n) \geq 0$, so $h(G) = 0$ for any goal G


Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

- **Theorem**: if h is admissible, A* search finds the optimal solution

<https://tutorcs.com>



Consistency

程序代写代做 CS编程辅导

- A heuristic is **consistent** if $h(n) \leq c(n, a, n') + h(n')$

- If h is consistent, then

- h is admissible, and

- $f(n)$ is nondecreasing along any path: $f(n') = g(n') + h(n')$


WeChat: cstutorcs

Assignment Project Exam Help
 $\geq g(n) + h(n)$
 $= f(n)$

- Consequently, when expanding a node, we cannot get a node with a smaller f , and so **the value of the best node on the frontier will never decrease.**

- When we dequeue a node labelled with a new state, we found the optimal path to that state:

- any other node m labelled with the same state satisfies $f(n) \leq f(m)$ and $h(n) = h(m)$, hence $g(n) \leq g(m)$.



Email: tutors@163.com

QQ: 749389476

<https://tutorcs.com>

Optimality of A* (based on consistency)

程序代写代做 CS编程辅导



Tutor CS

WeChat: cstutors

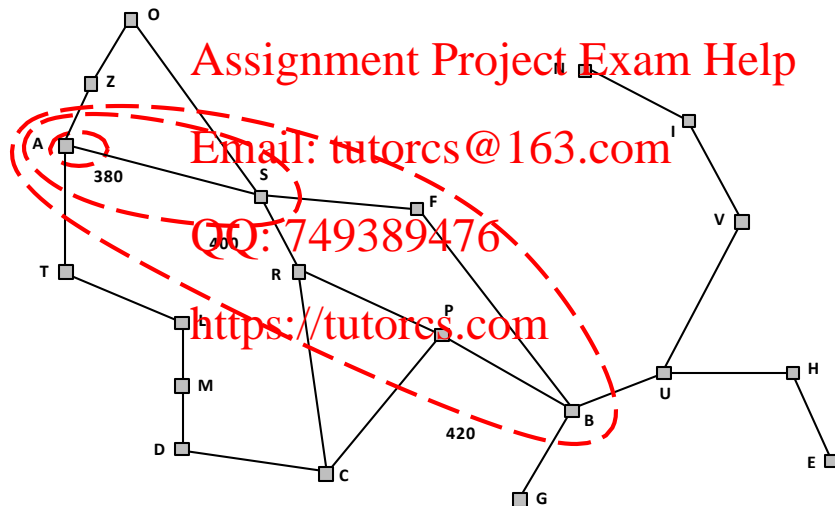
Assignment Project Exam Help

Email: tutorms@163.com

QQ: 749389476

https://tutorms.com

- **Consistency:** A* expands nodes in order of increasing f value
- Gradually expands “ f -contour” nodes
 - Contour i has all nodes with $f = f_i$, where $f_i < f_{i+1}$
 - (breadth-first expands layers, uniform-cost expands g -contours)



Properties of A*

程序代写代做CS编程辅导

- **A* expands**

- all nodes with $f(n) < C^*$
- some nodes with $f(n) = C^*$ (only one goal node)
- no nodes with $f(n) > C^*$



WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>


- Complete? Yes, unless there are infinitely many nodes with $f \leq C^*$

- Time? Exponential in [relative error in $h \times$ length of solution]

- Space? Exponential

- Optimal? Yes! Cannot expand f until f is finished

IDA* is a version of iterative deepening with a cutoff on f

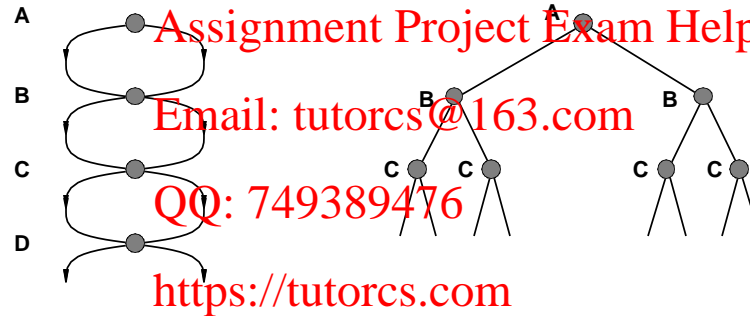


Tree Search and Repeated States

程序代写代做 CS编程辅导

- For many problems, the state space is a graph rather than a tree
- Cycles can prevent termination
- Failure to detect repeated states can turn a linear problem into an exponential one!

WeChat: cstutorcs



Graph search (simple)

程序代写代做 CS编程辅导

function GRAPH-SEARCH(*problem*, *frontier*) **returns** a solution, or failure

explored ← an empty set

frontier ← INSERT(M, INITIAL-STATE[*problem*]), *frontier*)

loop do

if *frontier* is empty **return** failure

node ← REMOVE-FRONT(*frontier*)

if GOAL-TEST(*problem*, STATE[*node*]) **then return** *node*

add *node* to *explored*

frontier ← INSERTNODES(EXPAND(*node*, *problem*), *frontier*)

WeChat: estutorcs

Assignment Project Exam Help

function INSERTNODES(*nodes*, *frontier*) **returns** updated frontier

for each *n* in *nodes* **do**

if $\nexists m$ in *explored* \cup *frontier* s.t. STATE[*n*] = STATE[*m*] **then**

add *n* to *frontier*

return *frontier*

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

- At most one instance of each state in $\text{explored} \cup \text{frontier}$
- All expanded nodes kept in memory!


Graph search and optimality

程序代写代做 CS编程辅导

- When seeking optimal solutions, multiple paths to the same state may need to be explored and compared to find the optimal



Number inside the nodes represents $g(n)$, i.e., the path cost



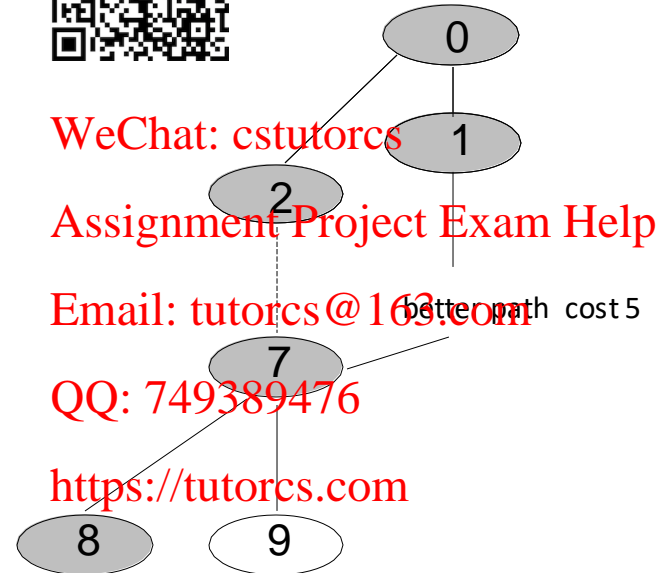
Graph search and optimality

程序代写代做 CS编程辅导

- When seeking optimal solution, multiple paths to the same state may need to be explored and compared to find the optimal



Number inside the nodes represents $g(n)$, i.e., the path cost




Graph search and optimality

程序代写代做 CS编程辅导

- We may need to keep the best node and update the depths and path-costs of the descendants of the first one.



Number inside the nodes represents $g(n)$, i.e., the path cost



Graph search and optimality

程序代写代做 CS编程辅导

- Trick to avoid updating descendants: re-open the explored node; its descendants will be updated if it is re-expanded.



Number inside the nodes represents $g(n)$, i.e., the path cost

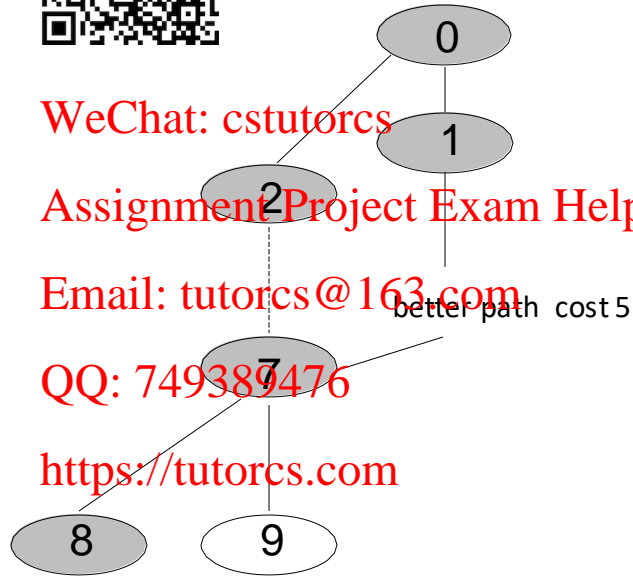
WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>



Graph search and optimality

程序代写代做 CS编程辅导

- Trick to avoid updating descendants: re-open the explored node; its descendants will be updated if it is re-expanded.



Number inside the nodes represents $g(n)$, i.e., the path cost



Graph search and optimality

程序代写代做 CS编程辅导

- Trick to avoid updating descendants: re-open the explored node; its descendants will be updated if it is re-expanded.



Number inside the nodes represents $g(n)$, i.e., the path cost



Graph Search (optimal)

程序代写代做 CS编程辅导

function INSERTNODES(*nodes*, *frontier*) **returns** updated frontier

for each *n* in *nodes* **do**

if $\nexists m$ in *explored* \cup

add *n* to *frontier*



STATE[*m*] = STATE[*n*] **then**

} Same as
before

else if PATH-COST[*n*] < PATH-COST[*m*]

PATH-COST[*m*] \leftarrow PATH-COST[*n*]

PARENT[*m*] \leftarrow PARENT[*n*]

ACTION[*m*] \leftarrow ACTION[*n*]

DEPTH[*m*] \leftarrow DEPTH[*n*]

if *m* in *explored* **then**

move *m* back to *frontier* // reopen *m*

return *frontier*

WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

} Found a
cheaper
path to *m*

- The internal block (**else if**) is needed with A*
 - *Uniform cost*: it also needs this block (unless all step costs are equal)
- **If *h* is consistent** (not just admissible), **no re-opening** (last 2 lines) is needed.
 - *Uniform cost*: no reopening is needed because $h = 0$ is consistent

<https://tutorcs.com>

Which algorithm and strategy to use

程序代写代做 CS编程辅导

| strategy | solution | useful when | algorithm | algorithm and state space |
|---------------|--------------------------|--------------------------------------|-------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| DFS | arbitrary | many solutions exist | DFS | <ul style="list-style-type: none"> tree-search for finite acyclic graphs recursive algorithm for finite acyclic graphs add cycle detection for finite graphs |
| BFS | shortest | shallow solutions exist | FIFO | <ul style="list-style-type: none"> tree search graph search (simple) may improve performance |
| UC | optimal | good admissible heuristic lacking | priority queue ordered by g | <ul style="list-style-type: none"> tree search for trees graph-search (simple) for equal step costs graph-search (optimal) for arbitrary step costs (no reopening needed) |
| A* | optimal | good admissible heuristics exist | priority queue ordered by f | <ul style="list-style-type: none"> tree search for trees and admissible heuristics graph-search (optimal) for admissible heuristics (no reopening needed for consistent heuristics) |
| greedy search | arbitrary but maybe good | good (inadmissible) heuristics exist | priority queue ordered by h | <ul style="list-style-type: none"> tree search for finite acyclic graphs graph search (simple) for finite graphs |



WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com


QQ: 749389476

https://tutorcs.com

Admissible heuristics – Example

程序代写代做 CS编程辅导

8-puzzle:



WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

- $h_1(n)$ = number of misplaced tiles
 - $h_1(\text{start state above}) = 6$
- $h_2(n)$ = total Manhattan distance
(i.e., no. of squares from desired location of each tile)
 - $h_2(\text{start state above}) = 4+0+3+3+1+0+2+1 = 14$

Dominance

程序代写代做 CS编程辅导

d : depth of the shallowest solution



- Given two admissible heuristics h_a and h_b , if $h_b(n) \geq h_a(n)$ for all n then
 - h_b dominates h_a and is better for search
 - In the 8-puzzle h_2 (Manhattan distance) dominates h_1 (# misplaced). Typical search costs:

| | | | |
|------------------------|-----------------------|---------------------------|------------------------------------|
| $d = 14$ | IDS = 3,473,941 nodes | $d = 24$ | IDS \approx 54,000,000,000 nodes |
| $A^*(h_1) = 539$ nodes | | $A^*(h_1) = 39,135$ nodes | |
| $A^*(h_2) = 113$ nodes | | $A^*(h_2) = 1,641$ nodes | |
- There is a trade-off between the accuracy of h and the time to compute h
- Given two admissible heuristics h_a and h_b
 - $h(n) = \max(h_a(n), h_b(n))$ is also admissible and dominates h_a and h_b
- Is $h_a(n) + h_b(n)$ admissible? No! It can double count work

WeChat: cstutorcs

Assignment Project Exam Help

Email: tutores@163.com

QQ: 749388476

https://tutores.com

Relaxed Problems

程序代写代做 CS编程辅导

- Admissible heuristics can be derived from the **optimal** solution cost of a **relaxed** version of the problem



- **Key point:** the optimal solution cost of a relaxed problem is no greater than the optimal solution cost of the real problem

WeChat: cstutorcs

- Rules of the 8-puzzle:

Assignment Project Exam Help

a tile can move from square A to square B if A is adjacent to B and B is blank; get all tiles in their correct positions.

Email: tutorcs@163.com

QQ: 749389476

- If we relax the rules so that a **tile can move anywhere**

- then $h_1(n)$ gives the shortest solution

<https://tutorcs.com>

- If we relax the rules so that a **tile can move to any adjacent square**

- then $h_2(n)$ gives the shortest solution

Relaxed Problems

程序代写代做 CS编程辅导

Rules of the 8-puzzle:

a tile can move from square A to square B if A is adjacent to B and B is blank; get all tiles into their correct positions.



- Relaxing the rules so that **only some tiles need to get in their correct positions** and solving the relaxed problem optimally yields another admissible heuristic



Start State Goal State


- Heuristics derived from the cost of an optimal solution to a smaller subproblem are used in **pattern databases** to store solutions for every possible subproblem up to a given size.

Relaxed Problems

程序代写代做 CS编程辅导

- Well-known example: **travelling salesperson problem** (TSP), i.e., find the least-cost tour visiting all cities exactly once

NP-Hard



- Minimum spanning tree** can be computed in $O(n^2)$
 - and the sum of the edge costs in an MST is a lower bound on the optimal (open) tour cost

QQ: 719389476

<https://tutorcs.com>

Announcements

程序代写代做 CS编程辅导



- **No class on Monday!** (Canberra Day)
- First Tutorials and Labs next Monday tutorial week after next.
- Quiz:
 - during your tutorial the tutor will give you a password for the quiz
 - After the tutorial, your quiz password will stop working
 - open book!
 - If you have an in-person tutorial: **you must bring a device** (e.g., laptop or tablet) to do the quiz. **Email: tutors@anu.edu.au**
- **Mid-semester exam will be on March 28**, 12pm-midnight on Wattle. (1.5h+15min reading time). It uses Proctorio, please make sure in advance that it works for you
- **Assignment 1 due March 31.** Due to the exam should we move it to a week later?

Assignment Project Exam Help

Email: tutors@anu.edu.au

QQ: 749389476

<https://tutorcs.com>



Australian
National
University

程序代写代做 CS编程辅导



Adversarial Search (game playing)

WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Chapter 5

Outline

程序代写代做 CS编程辅导



- Games
- Perfect play
 - minimax decisions
 - α - β pruning
- Imperfect decisions in real time

WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Adversarial Search Problems (Games)

程序代写代做 CS编程辅导



- Arise in competitive multi-agent environments
- In Game Theory, a multi-agent environment is called a game
- In AI, a game is often a deterministic, turn-taking, two-player, zero-sum game of perfect information:
 - deterministic
 - two agents
 - whose action alternate
 - utility values are opposite, e.g. $(+1, -1)$
 - fully observable
- We will write algorithms that play such games against an opponent.


WeChat: estutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>



Games Definition

程序代写代做 CS编程辅导

A Game consists of:

- sets of **players** P , **states** S (both of which a player to play), and **moves** M
- an **initial state** $s_0 \in S$ which defines how the game is set up
- $\text{Player}(s) \in P$: defines **the player to move** in state s
- $\text{Moves}(s) \in 2^M$: defines **the set of legal moves** in state s
- $\text{Result}(s, m) \in S$: defines the **result** of performing move m in state s
- $\text{Terminal}(s) \in \{T, F\}$: the **terminal test** says whether the game is over
- $\text{Utility}(s, p) \in \mathbb{R}$: the **utility function** gives a numeric value to terminal states from the point of view of a given player, e.g.:
 - $\{+1, -1, 0\}$ for chess
 - $\{-192, \dots, 192\}$ for backgammon



WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 74938476

<https://tutorcs.com>

Strategy

程序代写代做 CS 编程辅导

- “Unpredictable” opponent
⇒ **solution** for a player is not a sequence of actions but a **strategy**
- A **strategy** for a player: a function mapping the player's states to (legal) moves.
- A **winning strategy** always lead the player to a win from so



WeChat: cstutorcs

Assignment Project Exam Help


Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Example: Tic-tac-toe

程序代写代做 CS编程辅导



WeChat: cstutores

- States? content of each cell {X, O, empty}, player to play
- Moves? an empty cell
- Result? content of chosen cell is X or O depending on the player playing;
next player
- Terminal test? are 3 Os or 3 Xs aligned or is the board full?
- Utility function? for a given player gives +1 if player has aligned 3 tokens, -1 if his opponent has, and 0 otherwise. There is a draw strategy.

Assignment Project Exam Help

Email: tutorcs@163.com

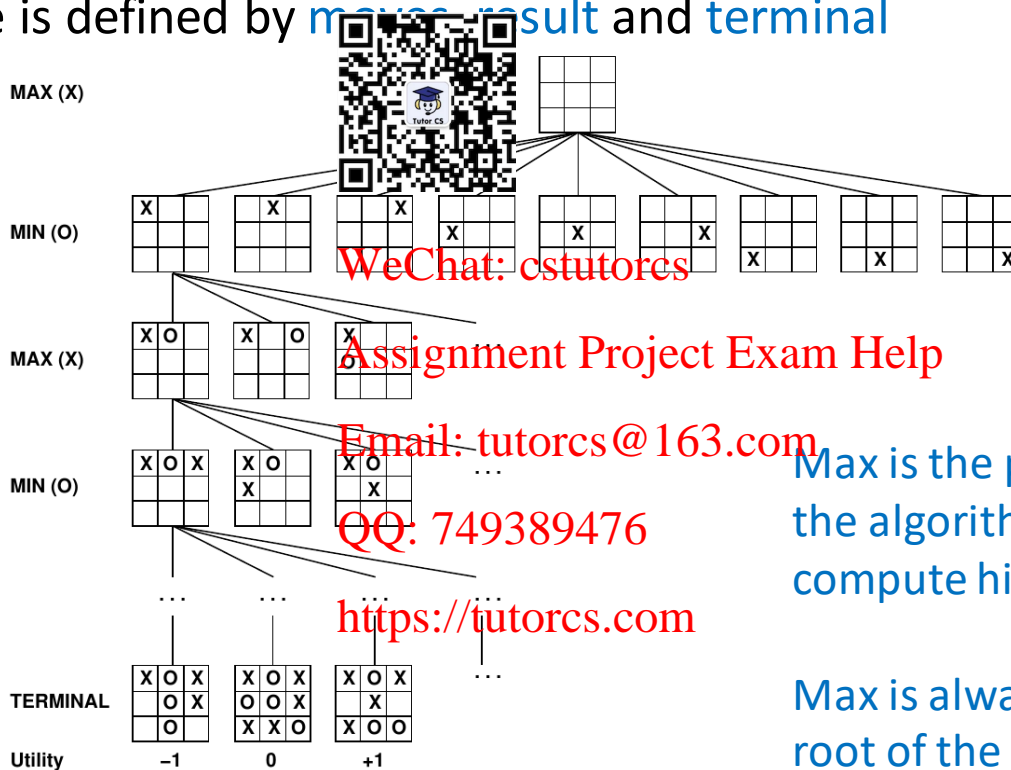
QQ: 749389476

<https://tutorcs.com>

Game tree (2-player, deterministic, turns)

程序代写代做 CS编程辅导

The game tree is defined by **moves**, **result** and **terminal**



Max is the player using
the algorithm to
compute his strategy

Max is always at the
root of the tree

Minimax

程序代写代做 CS编程辅导

- **Perfect play** for deterministic, two-player, zero-sum, perfect-information games
- **Idea:** choose move to position with highest minimax value
 - best achievable utility against best possible opponent




$$\text{MINIMAX-VALUE}(s) = \begin{cases} \text{UTILITY}(s, \text{MAX}) & \text{if } \text{TERMINAL}(s) \\ \max_{m \in \text{MOVES}(s)} \text{MINIMAX-VALUE}(\text{RESULT}(s, m)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{m \in \text{MOVES}(s)} \text{MINIMAX-VALUE}(\text{RESULT}(s, m)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$

WeChat: cstutores

Assignment Project Exam Help

E.g. 2-ply game: MAX



Email: tutores@163.com

QQ: 749389476

<https://tutores.com>

Minimax

程序代写代做 CS编程辅导



WeChat: tutorcs

Assignment Project Exam Help

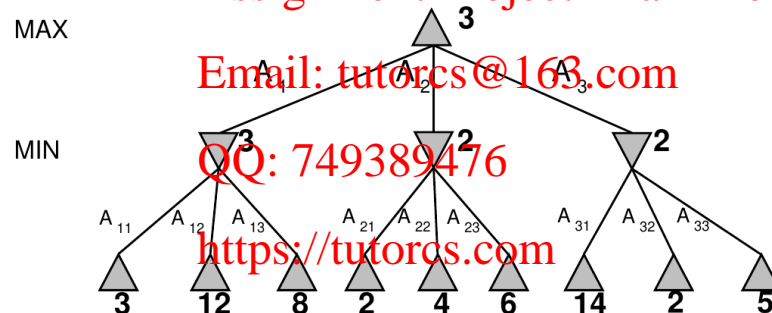
Email: tutorcs@163.com

QQ: 749389476

https://tutorcs.com

Computing the Minimax value:

- Apply utility function to each leaf node of the game tree
- Back-up values from the leaf nodes through inner nodes up to the root:
 - **min** node: compute the min of its children values
 - **max** node: compute the max of its children values
- At the root: choose the move leading to the child of highest value



- Better method: use a depth-first like approach to save space

Minimax Algorithm


程序代写代做 CS编程辅导

function MINIMAX-DECISION(*state*) **returns** a move

inputs: *state*, current player

$v \leftarrow \text{MAX-VALUE}(state, \text{current player})$

return the move *m* such that $U(state, m) = v$



function MAX-VALUE(*state*) **returns** a utility value

if TERMINAL(*state*) **then return** UTILITY(*state*, MAX)

$v \leftarrow -\infty$

for *m* in MOVES(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(state, m)))$

return *v*

function MIN-VALUE(*state*) **returns** a utility value

if TERMINAL(*state*) **then return** UTILITY(*state*, MAX)

$v \leftarrow +\infty$

for *m* in MOVES(*state*) **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(state, m)))$

return *v*

WeChat: estutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

Properties of Minimax

程序代写代做 CS编程辅导



WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

- Complete? Yes, if tree is finite
- Optimal? Yes, against an optimal opponent. Otherwise?
- Time complexity? $O(b^m)$
- Space complexity? $O(bm)$ (depth-first exploration)
 - For chess, $b \approx 35$, $m \approx 100$ for “reasonable” games
 \Rightarrow exact solution completely infeasible
 - But do we need to explore every path?

b : max. branching factor
 d : depth of the shallowest solution
 m : max. depth of the state space

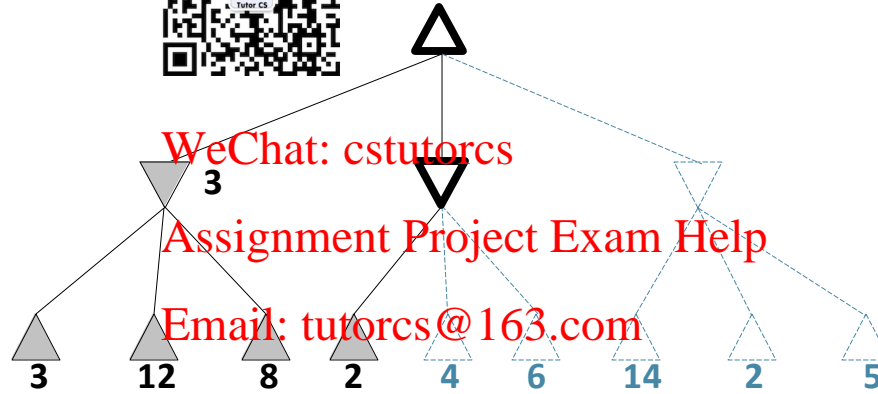
Do we need to explore every path?

程序代写代做 CS编程辅导



MAX

MIN



WeChat: cstutorcs

Assignment Project Exam Help


Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

α - β pruning

程序代写代做 CS编程辅导



WeChat: cstutorcs

MAX
Assignment Project Exam Help

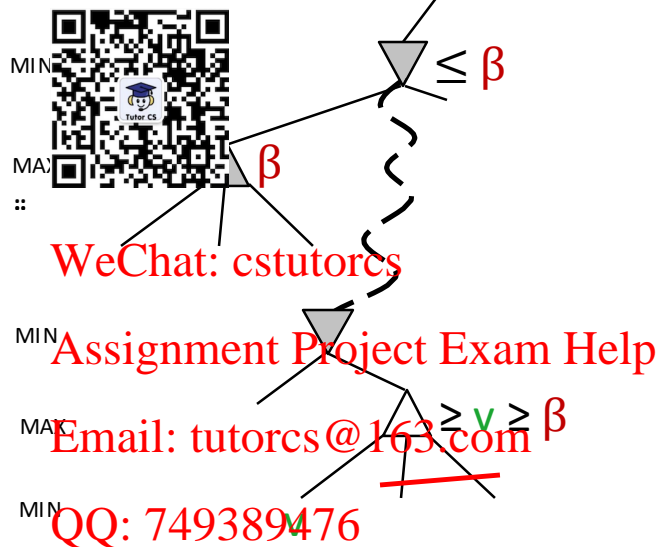
MIN
Email: tutorcs@163.com

MAX
QQ: 749389476

<https://tutorcs.com>

- α is the best value (to max, i.e. highest) found so far
- If v is not better (greater) than α , max will avoid it
 - prune that branch
- Define β similarly for min

程序代写代做 CS编程辅导



- β is the best value (to min, i.e. lowest) found so far
- If v is not worse (lower) than β , min will avoid it
 - prune that branch

α - β pruning

程序代写代做 CS编程辅导

- Idea:

- α is the best (largest) value for \max on path to current node
- β is the best (lowest) value for \min on path to current node
- **Some values** outside the interval $[\alpha, \beta]$ can be pruned

WeChat: cstutorcs

- Algorithm:

- Node passes its current values for α and β to its children in turn
- Child passes back up its value v to Node
- Node updates its current value v (max or min with child's value)
- Node checks whether $v \leq \alpha$ (for min node) or $v \geq \beta$ (for max node)
 - **If so**, child's siblings can be pruned and v returned to Parent
 - **Otherwise** β (min) or α (max) is updated

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

https://tutorcs.com

The α - β algorithm

程序代写代做 CS编程辅导

function ALPHA-BETA-DECISION(*state*) **returns** a move
 $v \leftarrow \text{MAX-VALUE}(\text{state})$
return the move *m* in *state* with value *v*

function MAX-VALUE(*state*) **returns** a utility value
inputs: *state*, current state in game
 α , the value of the best choice for MAX so far
 β , the value of the best choice for MIN so far

if TERMINAL(*state*) **then return** UTILITY(*state*, MAX)
 $v \leftarrow -\infty$
for *m* in MOVES(*state*) **do**
 $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(m, s), \alpha, \beta))$
if $v \geq \beta$ **then return** *v*
 $\alpha \leftarrow \text{MAX}(\alpha, v)$
return *v*

function MIN-VALUE(*state*, α , β) **returns** a utility value
 same as MAX-VALUE but with roles of α , β reversed

WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>